

# Car Pricing

2022-08-18

## R Markdown

```
#Removing from memory, packages
rm(list=ls())
library(Matrix) # For matrix computations
library(olsrr) # For Variable selection algorithms

## Warning: package 'olsrr' was built under R version 4.2.1

##
## Attaching package: 'olsrr'

## The following object is masked from 'package:datasets':
##
##      rivers

library(car) # For VIF

## Loading required package: carData

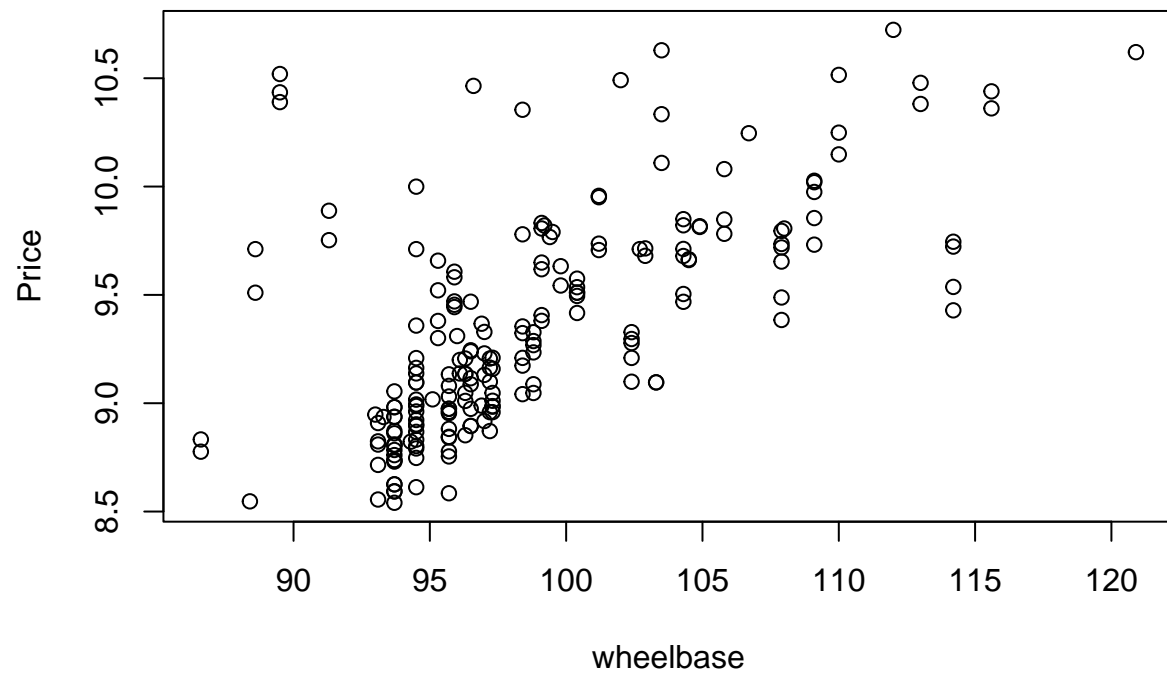
library(mctest) # For variance decomposition algorithms


#reading data
library(readxl)
data=read_excel("C:/Users/saubh/Downloads/CarPrice_Assignment.xlsx",col_names=TRUE)
data=data.frame(data)
attach(data)
n=length(price)


#define response, predictors, design matrix, cs matrix
y=data[,14]
y=log(y) ##### Idea of Box cox method was used
data_pred = data[,-c(14)]
p=length(data_pred[1,]) #equals 13 (variable selection needed)
X = cbind(rep(1,n),as.matrix(data_pred))
Xcs = data.frame(sqrt(1/(n-1))*scale.default(data_pred, center=TRUE,scale=TRUE))
Xcs=as.matrix(Xcs)
det(t(Xcs)%*%Xcs) # equals 3.3e-07 (suspecting multicollinearity)
```


```
## [1] 3.349118e-07
```

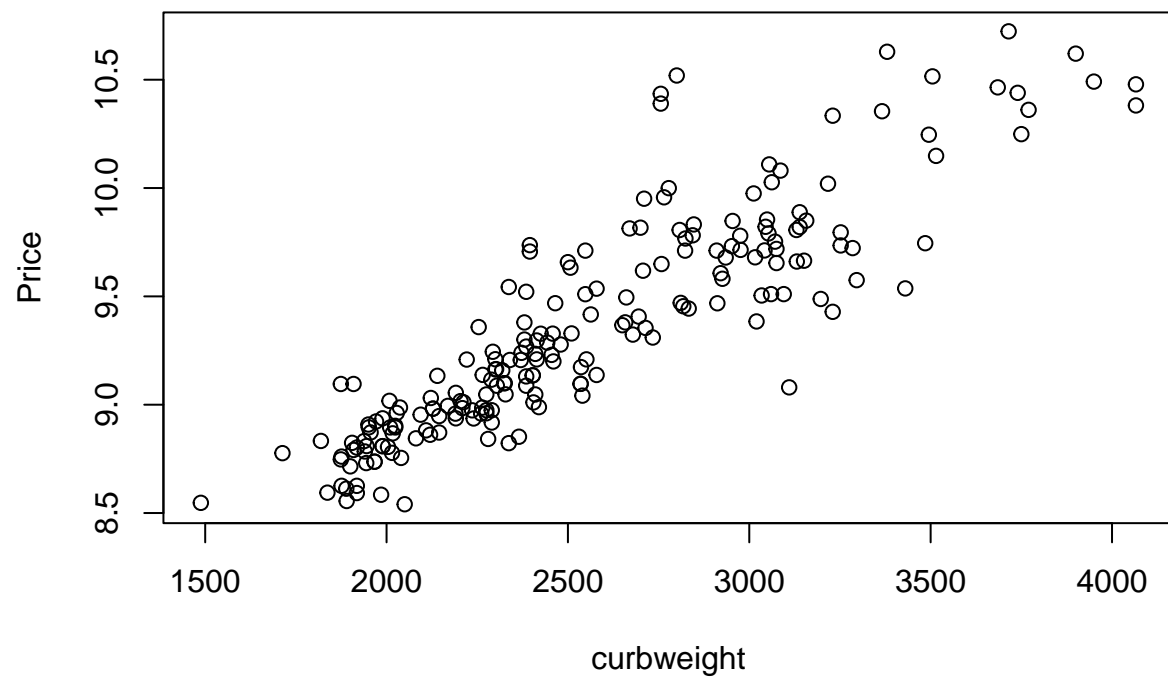
```
#pairwise scatterplot  
#par(mfrow=c(3,5))  
for(i in 1:12){  
  plot(data[,i],y,xlab=names(data)[i],ylab="Price")  
}
```





















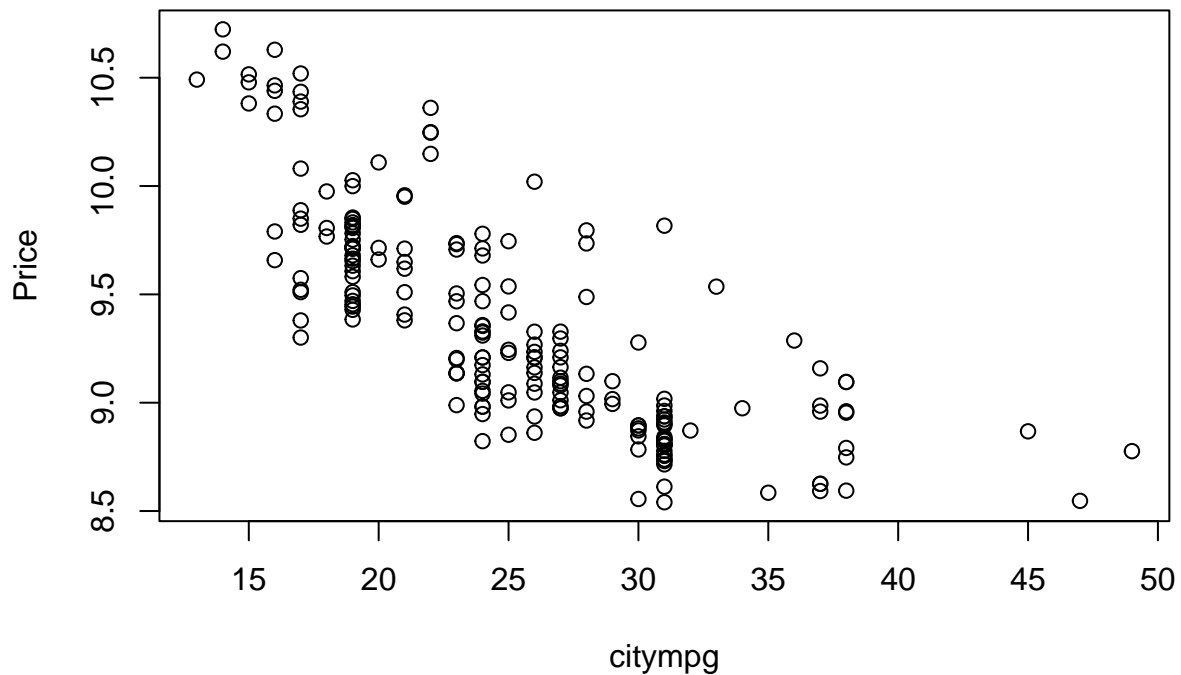












```
#fitting
fit=lm(y~.,data=data_pred)
summary(fit) # r2=0.8792
```

```
##
## Call:
## lm(formula = y ~ ., data = data_pred)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.58368 -0.10611 -0.02615  0.10144  0.44373
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   6.024e+00  8.663e-01   6.953 5.51e-11 ***
## wheelbase     3.222e-03  5.700e-03   0.565 0.572591
## carlength     6.875e-04  3.152e-03   0.218 0.827575
## carwidth      2.036e-02  1.396e-02   1.458 0.146384
## carheight     4.066e-03  7.700e-03   0.528 0.598093
## curbweight    2.708e-04  9.857e-05   2.747 0.006585 **
## enginesize     2.679e-03  7.851e-04   3.413 0.000785 ***
## boreratio     -1.734e-02  6.785e-02  -0.256 0.798530
## stroke        -1.155e-01  4.418e-02  -2.615 0.009648 **
## compressionratio 2.262e-02  4.704e-03   4.808 3.08e-06 ***
## horsepower     3.050e-03  9.201e-04   3.315 0.001096 **
## peakrpm        9.387e-05  3.806e-05   2.466 0.014541 *
```

```
## citympg          -3.170e-02  1.009e-02  -3.143 0.001938 **
## highwaympg      1.645e-02  9.064e-03   1.815 0.071063 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1809 on 191 degrees of freedom
## Multiple R-squared:  0.8792, Adjusted R-squared:  0.871
## F-statistic: 107 on 13 and 191 DF, p-value: < 2.2e-16
```

```
#####
```

```
# Evaluating all possible subset linear regression models is not advisable since total no. would be 213
```

```
# Forward selection method
```

```
f = ols_step_forward_p(fit, details = FALSE)
summary(f$model) # r2=0.8789
```

```
##
## Call:
## lm(formula = paste(response, "~", paste(preds, collapse = " + ")),
##     data = l)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.58880 -0.10819 -0.02394  0.10387  0.45408
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   6.162e+00  7.455e-01   8.266 2.15e-14 ***
## curbweight    2.837e-04  9.364e-05   3.029 0.002784 **
## horsepower    2.936e-03  8.883e-04   3.305 0.001130 **
## compressionratio 2.266e-02  4.667e-03   4.855 2.47e-06 ***
## citympg      -3.223e-02  9.483e-03  -3.399 0.000820 ***
## carwidth      1.940e-02  1.339e-02   1.448 0.149195
## enginesize    2.621e-03  7.709e-04   3.400 0.000817 ***
## stroke       -1.182e-01  4.242e-02  -2.787 0.005849 **
## peakrpm       9.564e-05  3.633e-05   2.633 0.009154 **
## highwaympg    1.705e-02  8.800e-03   1.937 0.054191 .
## wheelbase     5.110e-03  4.516e-03   1.131 0.259292
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1798 on 194 degrees of freedom
## Multiple R-squared:  0.8789, Adjusted R-squared:  0.8727
## F-statistic: 140.8 on 10 and 194 DF, p-value: < 2.2e-16
```

```
# backward elimination method
```

```
b = ols_step_backward_p(fit, details=FALSE)
summary(b$model) # r2=0.8789
```

```
##
## Call:
```

```
## lm(formula = paste(response, "~", paste(preds, collapse = " + ")),
##     data = l)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.58880 -0.10819 -0.02394  0.10387  0.45408
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    6.162e+00  7.455e-01   8.266 2.15e-14 ***
## wheelbase      5.110e-03  4.516e-03   1.131 0.259292
## carwidth       1.940e-02  1.339e-02   1.448 0.149195
## curbweight     2.837e-04  9.364e-05   3.029 0.002784 **
## enginesize     2.621e-03  7.709e-04   3.400 0.000817 ***
## stroke        -1.182e-01  4.242e-02  -2.787 0.005849 **
## compressionratio 2.266e-02  4.667e-03   4.855 2.47e-06 ***
## horsepower     2.936e-03  8.883e-04   3.305 0.001130 **
## peakrpm        9.564e-05  3.633e-05   2.633 0.009154 **
## citympg       -3.223e-02  9.483e-03  -3.399 0.000820 ***
## highwaympg     1.705e-02  8.800e-03   1.937 0.054191 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1798 on 194 degrees of freedom
## Multiple R-squared:  0.8789, Adjusted R-squared:  0.8727
## F-statistic: 140.8 on 10 and 194 DF,  p-value: < 2.2e-16
```

```
# Step-wise selection method
s = ols_step_both_p(fit, details=FALSE)
summary(s$model) # r2=0.8781
```

```
##
## Call:
## lm(formula = paste(response, "~", paste(preds, collapse = " + ")),
##     data = l)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.60078 -0.11198 -0.02174  0.09495  0.45516
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    6.138e+00  7.457e-01   8.232 2.61e-14 ***
## curbweight     3.308e-04  8.395e-05   3.940 0.000113 ***
## horsepower     2.538e-03  8.160e-04   3.110 0.002153 **
## compressionratio 2.195e-02  4.628e-03   4.743 4.06e-06 ***
## citympg       -3.091e-02  9.416e-03  -3.282 0.001221 **
## carwidth       2.637e-02  1.190e-02   2.216 0.027862 *
## enginesize     2.598e-03  7.711e-04   3.368 0.000911 ***
## stroke        -1.179e-01  4.245e-02  -2.777 0.006014 **
## peakrpm        9.546e-05  3.635e-05   2.626 0.009324 **
## highwaympg     1.595e-02  8.753e-03   1.822 0.069985 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.1799 on 195 degrees of freedom
## Multiple R-squared:  0.8781, Adjusted R-squared:  0.8725
## F-statistic: 156.1 on 9 and 195 DF,  p-value: < 2.2e-16

# g = r2(M)/r2(whole)
# So for forward/backward selection, g is close to 1

#####

# we drop the boreratio, carlength, carheight, updating X, Xcs
data_pred2=subset(data_pred, select=-c(carlength, boreratio, carheight))
X = cbind(rep(1, n), as.matrix(data_pred2))
Xcs = data.frame(sqrt(1/(n-1))*scale.default(data_pred2, center=TRUE, scale=TRUE))
Xcs=as.matrix(Xcs)
yxs=(y-mean(y))/((n-1)*(sd(y)))

# fitting MLRM with CS data
scaled_fit = lm(yxs~.-1, data=data.frame(Xcs))
summary(scaled_fit)

##
## Call:
## lm(formula = yxs ~ . - 1, data = data.frame(Xcs))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.005729 -0.001053 -0.000233  0.001011  0.004418
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## wheelbase      0.004276   0.003770   1.134  0.258065
## carwidth       0.005783   0.003983   1.452  0.148148
## curbweight     0.020528   0.006759   3.037  0.002714 **
## enginesize     0.015168   0.004450   3.409  0.000792 ***
## stroke        -0.005152   0.001844  -2.794  0.005723 **
## compressionratio 0.012507   0.002569   4.867  2.33e-06 ***
## horsepower     0.016135   0.004869   3.314  0.001097 **
## peakrpm        0.006340   0.002402   2.639  0.008976 **
## citympg       -0.029306   0.008599  -3.408  0.000795 ***
## highwaympg      0.016314   0.008400   1.942  0.053573 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.001745 on 195 degrees of freedom
## Multiple R-squared:  0.8789, Adjusted R-squared:  0.8727
## F-statistic: 141.6 on 10 and 195 DF,  p-value: < 2.2e-16

# correlation matrix
corr_mat=t(Xcs)%*(Xcs)
corr_mat # high pairwise correlation between predictors is present (suspecting multicollinearity)
```



```
##          wheelbase  carwidth  curbweight  enginesize      stroke
## wheelbase      1.0000000  0.7951436  0.7763863  0.56932868  0.16095905
## carwidth       0.7951436  1.0000000  0.8670325  0.73543340  0.18294169
## curbweight     0.7763863  0.8670325  1.0000000  0.85059407  0.16879004
## enginesize     0.5693287  0.7354334  0.8505941  1.00000000  0.20312859
## stroke        0.1609590  0.1829417  0.1687900  0.20312859  1.00000000
## compressionratio 0.2497858  0.1811286  0.1513617  0.02897136  0.18611011
## horsepower     0.3532945  0.6407321  0.7507393  0.80976865  0.08093954
## peakrpm       -0.3604687 -0.2200123 -0.2662432 -0.24465983 -0.06796375
## citympg       -0.4704136 -0.6427043 -0.7574138 -0.65365792 -0.04214475
## highwaympg    -0.5440819 -0.6772179 -0.7974648 -0.67746991 -0.04393093
##          compressionratio  horsepower      peakrpm      citympg
## wheelbase      0.24978585  0.35329448 -0.36046875 -0.47041361
## carwidth       0.18112863  0.64073208 -0.22001230 -0.64270434
## curbweight     0.15136174  0.75073925 -0.26624318 -0.75741378
## enginesize     0.02897136  0.80976865 -0.24465983 -0.65365792
## stroke        0.18611011  0.08093954 -0.06796375 -0.04214475
## compressionratio 1.00000000 -0.20432623 -0.43574051  0.32470142
## horsepower     -0.20432623  1.00000000  0.13107251 -0.80145618
## peakrpm       -0.43574051  0.13107251  1.00000000 -0.11354438
## citympg       0.32470142 -0.80145618 -0.11354438  1.00000000
## highwaympg    0.26520139 -0.77054389 -0.05427481  0.97133704
##          highwaympg
## wheelbase      -0.54408192
## carwidth       -0.67721792
## curbweight     -0.79746479
## enginesize     -0.67746991
## stroke        -0.04393093
## compressionratio 0.26520139
## horsepower     -0.77054389
## peakrpm       -0.05427481
## citympg       0.97133704
## highwaympg    1.00000000
```

```
# Determinant of X'X
det(t(Xcs)%*%Xcs) # equals 1.35e-05 (suspecting multicollinearity)
```

```
## [1] 1.356289e-05
```

```
# Computing VIF
round(vif(scaled_fit),2)
```

```
## Warning in vif.default(scaled_fit): No intercept: vifs may not be sensible.
```

```
##          wheelbase      carwidth      curbweight      enginesize
##          4.67          5.21          15.01          6.51
##          stroke compressionratio      horsepower      peakrpm
##          1.12          2.17          7.79          1.90
##          citympg      highwaympg
##          24.30          23.19
```

```

# We use cut-off 10.
# Multi-collinearity is suspected and the regression coefficients corresponding to curbweight, highwaympg

# Condition indices
e = eigen(corr_mat)$values # eigen(A) computes both eigenvalues and eigenvectors
Condition_indices = array(0)
for(i in 1:length(e))
{
  Condition_indices[i] = max(e)/e[i]
}
round(Condition_indices,2)

```

```
## [1] 1.00 2.67 5.56 9.37 10.11 17.49 35.95 62.96 101.24 243.46
```

```

# We use cut-off 25.
# Last four condition indices are larger than 25.
# Last four principal components of the predictor variables are suspected to be responsible for multi-c

# Calculating measures based on variance decomposition
scaled_fit_intercept = lm(y ~ ., data = data.frame(Xcs))
VP = eigprop(scaled_fit_intercept, Inter=FALSE)
VP$pi[c(7,9,8,3,10),]

```

```

##      wheelbase      carwidth  curbweight  enginesize      stroke
## 7  0.258547720 0.7955971264 0.0126775826 0.1237058744 0.003380628
## 9  0.285083125 0.0045815702 0.9743783365 0.1363310742 0.001059359
## 8  0.215837989 0.1044133174 0.0019231834 0.5510365296 0.024772665
## 3  0.003845473 0.0001323175 0.0001674864 0.0003940061 0.813182732
## 10 0.009830957 0.0072100410 0.0071348413 0.0498118113 0.011116077
##      compressionratio  horsepower    peakrpm      citympg  highwaympg
## 7      0.0215145347 0.0065624467 0.12127661 8.145205e-04 0.0017293219
## 9      0.2463499096 0.0733056838 0.03420179 2.559639e-02 0.0140254593
## 8      0.0051883280 0.7369467746 0.16912461 2.855819e-04 0.0230609106
## 3      0.0001899242 0.0009064886 0.05235804 1.716118e-05 0.0001192479
## 10     0.0087898641 0.1089218345 0.01131816 9.387845e-01 0.9241836241

```

```

# Regression coefficient correspond to :
# carwidth is suffering from the 7th pc. The measure is 0.79
# curbweight is suffering from the 9th pc. The measure is 0.97
# stroke is suffering from the 3rd pc. The measure is 0.81
# horsepower is suffering from the 8th pc. The measure is 0.73
# citympg is suffering from the 10th pc. The measure is 0.93
# highwaympg is suffering from the 10th pc. The measure is 0.92

```

```

##### PC REGRESSION #####
pca = prcomp(Xcs)
pca

```

```
## Standard deviations (1, ..., p=10):
## [1] 0.16168414 0.09898412 0.06858343 0.05280622 0.05084676 0.03865789
## [7] 0.02696705 0.02037598 0.01606878 0.01036220
##
## Rotation (n x k) = (10 x 10):
##
##          PC1          PC2          PC3          PC4          PC5
## wheelbase    0.321675796 -0.29350643  0.13125872 -0.266763140 -0.53481463
## carwidth      0.382912594 -0.16192752  0.02572578 -0.262279485 -0.09225934
## curbweight    0.415504492 -0.11917170  0.04911328 -0.049982240  0.05467128
## enginesize    0.381042448 -0.05618778 -0.04959359  0.282978977  0.40397773
## stroke        0.077279467 -0.21532032 -0.93369364  0.179148820 -0.19020244
## compressionratio -0.007058283 -0.59477638 -0.01988290 -0.483030223  0.43602487
## horsepower    0.363671300  0.23106403 -0.08231035  0.058430375  0.49907480
## peakrpm       -0.068098222  0.53576773 -0.30859438 -0.712656903  0.08151663
## citympg       -0.374996408 -0.28060028 -0.02000226 -0.027193513  0.14177905
## highwaympg    -0.385391845 -0.22578428 -0.05150709 -0.005061678  0.19276413
##
##          PC6          PC7          PC8          PC9          PC10
## wheelbase   -0.31684921  0.42319237 -0.29215713 -0.264790636 -0.03170910
## carwidth    -0.29288094 -0.78436962  0.21470254 -0.035467512  0.02869205
## curbweight   0.04373586  0.16801249 -0.04944459  0.877680677 -0.04843218
## enginesize  -0.35929495  0.34552816  0.55101518 -0.216139982 -0.08425060
## stroke       0.06611404 -0.02367136 -0.04841679  0.007895789  0.01649373
## compressionratio 0.43118054  0.08320891  0.03087469 -0.167775849 -0.02043683
## horsepower  -0.11262205 -0.08708039 -0.69725336 -0.173422599  0.13632123
## peakrpm     -0.18281219  0.18467115  0.16477792  0.058436518 -0.02167789
## citympg     -0.47275313  0.05418397 -0.02424210  0.180991663  0.70684031
## highwaympg  -0.47215927 -0.07712475 -0.21280370  0.130877143 -0.68509935
```

*# prcomp(A) performs principal component analysis on the data matrix A.*

```
pc_variances = summary(pca) # summary(prcomp()) prints sample variance of the princpal components.
pc_variances
```

```
## Importance of components:
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
## Standard deviation    0.1617 0.09898 0.06858 0.05281 0.05085 0.03866 0.02697
## Proportion of Variance 0.5333 0.19988 0.09596 0.05689 0.05274 0.03049 0.01484
## Cumulative Proportion 0.5333 0.73317 0.82912 0.88601 0.93875 0.96924 0.98407
##
##          PC8          PC9          PC10
## Standard deviation    0.02038 0.01607 0.01036
## Proportion of Variance 0.00847 0.00527 0.00219
## Cumulative Proportion 0.99254 0.99781 1.00000
```

```
pc = pca$x # The i-th column of prcomp(A)$x provides the i-th principal component of the data matrix A.
pc_data = data.frame(pc)
summary(lm(ycs~.-1,pc_data)) # Fit MLRM for the standarized response on the principal components obtain
```

```
##
```

```
## Call:
```

```
## lm(formula = ycs ~ . - 1, data = pc_data)
```

```
##
```

```
## Residuals:
```

```
##          Min          1Q          Median          3Q          Max
```

```
## -0.005729 -0.001053 -0.000233 0.001011 0.004418
##
## Coefficients:
##      Estimate Std. Error t value Pr(>|t|)
## PC1    0.0275510  0.0007554  36.470 < 2e-16 ***
## PC2   -0.0001546  0.0012340  -0.125  0.90043
## PC3    0.0019895  0.0017809   1.117  0.26531
## PC4   -0.0092159  0.0023130  -3.984 9.55e-05 ***
## PC5    0.0184219  0.0024022   7.669 8.04e-13 ***
## PC6    0.0006269  0.0031596   0.198  0.84292
## PC7    0.0040460  0.0045293   0.893  0.37280
## PC8   -0.0049957  0.0059944  -0.833  0.40565
## PC9    0.0056652  0.0076012   0.745  0.45699
## PC10  -0.0324110  0.0117873  -2.750  0.00653 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.001745 on 195 degrees of freedom
## Multiple R-squared:  0.8789, Adjusted R-squared:  0.8727
## F-statistic: 141.6 on 10 and 195 DF,  p-value: < 2.2e-16

fit_good=lm(yes~.-1,pc_data[,c(1,4,5,10)])
summary(fit_good) # Fit final multiple linear regression model with significant principal components.

##
## Call:
## lm(formula = yes ~ . - 1, data = pc_data[, c(1, 4, 5, 10)])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0057926 -0.0011610 -0.0002189  0.0009591  0.0046280
##
## Coefficients:
##      Estimate Std. Error t value Pr(>|t|)
## PC1    0.0275510  0.0007504  36.713 < 2e-16 ***
## PC4   -0.0092159  0.0022977  -4.011 8.53e-05 ***
## PC5    0.0184219  0.0023863   7.720 5.36e-13 ***
## PC10  -0.0324110  0.0117094  -2.768  0.00617 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.001733 on 201 degrees of freedom
## Multiple R-squared:  0.8769, Adjusted R-squared:  0.8744
## F-statistic: 357.8 on 4 and 201 DF,  p-value: < 2.2e-16

round(pca$rotation[,c(1,4,5,10)],2) # This provides the linear combination for the significant principal components


##           PC1    PC4    PC5    PC10
## wheelbase  0.32 -0.27 -0.53 -0.03
## carwidth   0.38 -0.26 -0.09  0.03
## curbweight  0.42 -0.05  0.05 -0.05
## enginesize  0.38  0.28  0.40 -0.08
## stroke     0.08  0.18 -0.19  0.02
```

```
## compressionratio -0.01 -0.48 0.44 -0.02
## horsepower      0.36 0.06 0.50 0.14
## peakrpm         -0.07 -0.71 0.08 -0.02
## citympg         -0.37 -0.03 0.14 0.71
## highwaympg      -0.39 -0.01 0.19 -0.69
```

```
##### Assumption Verification #####
```

```
# Normality assumption
```

```
ols_plot_resid_qq(fit_good) # Outliers present and hence we can note departure from normality
```




```
ols_test_normality(fit_good) # All tests result in rejection of null hypothesis; therefore normality as
```

```
## Warning in ks.test.default(y, "pnorm", mean(y), sd(y)): ties should not be
## present for the Kolmogorov-Smirnov test
```

```
## -----
##      Test           Statistic      pvalue
## -----
## Shapiro-Wilk        0.9746        9e-04
## Kolmogorov-Smirnov   0.0601        0.4487
## Cramer-von Mises     68.0947        0.0000
## Anderson-Darling     1.6648        3e-04
## -----
```

```
# Homoscedasticity assumption
```


```
ols_plot_resid_fit(fit_good) # Clustering is observed and hence homoscedasticity assumption does not hold
```



```
# Assumption of random errors being uncorrelated
```

```
acf(fit_good$residuals, plot=TRUE) # Autocorrelation is present
```


### Series fit\_good\$residuals



##### Outlier detection #####

```
lev1=ols_plot_resid_lev(fit_good) # Computing leverage measure for all observations
```

## Outlier and Leverage Diagnostics for ycs



```
# leverage points
```

```
which(lev1$data$color=="leverage")
```

```
## [1] 9 48 49 51 73 74 75 91 106 110 114 127 145 150 160 166 167 183 185
```

```
## [20] 204
```

```
# outliers
```

```
which(lev1$data$color=="outlier")
```

```
## [1] 11 12 13 14 15 17 45 46 67 102 103 104
```


```
# outlier and leverage
```

```
which(lev1$data$color=="outlier & leverage")
```

```
## [1] 19 31 50 128 129 130
```

```
lev2=ols_plot_cooks_d_bar(fit_good) # Computing Cook's distance statistics
```



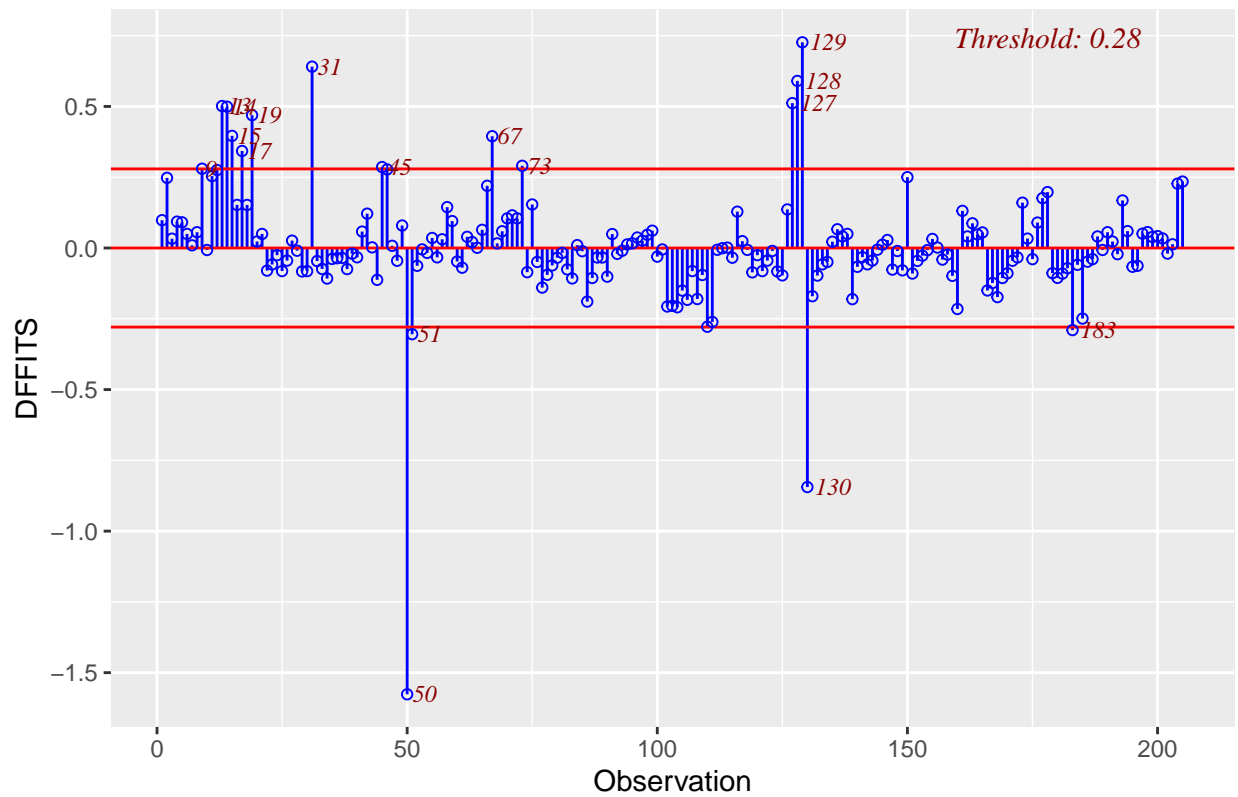


```
which(lev2$data$color=="outlier")
```

```
## [1] 9 13 14 15 17 19 31 45 50 51 67 73 127 128 129 130 183
```

```
lev3=ols_plot_dffits(fit_good) # Computing DFFITS statistics
```

## Influence Diagnostics for ycs



```
which(lev3$data$color=="outlier")
```

```
## [1] 9 13 14 15 17 19 31 45 50 51 67 73 127 128 129 130 183
```

```
ols_plot_dfbetas(fit_good)# Computing DFBETAS statistics
```

